NotebookLM

Exported on: 5/16/2026, 9:41:38 AM

Topic	Domain	Key Findings	Performance Impact Factors	Governance and Risks	Recommended Strategies	Source
Context Window and Positioning (Lost in the Middle)	LLM Technical Limitations / NLP	Models consistently underutilize information in the middle of long contexts, resulting in a U-shaped performance curve. Accuracy can drop by 20%+ compared to information at the beginning or end.	Positional bias (primacy and recency), context length (4k to 2M+), distractor volume, and attention dilution.	Risks of "context rot" and superlinear cost growth. Overfilling context leads to performance underperforming closed-book baselines.	Place critical info at the beginning or end; use RAG to ground responses; implement re-ranking to prioritize evidence; use structured attention anchors.	[1-7]
Agentic AI & Multi-Agent Systems (MAS)	AI Agents / Enterprise AI	Shift from isolated agents to autonomous collectives that plan, execute, and iterate on tasks. 75% of firms building in-house architectures are predicted to fail without external platforms.	Goal-directed planning, tool utilization, communication overhead, and the "complexity ceiling" (coordination failures often occur beyond 5 agents).	Accountability for autonomous logic; "black box" risks; risk of cascading failures (seen in 21% of enterprises); potential for improper delegation of professional responsibility.	Use a "human-in-the-loop" approach; implement failure attribution mechanisms; adopt "Agent Operating Systems" for structure; verify substantive conclusions at checkpoints.	[1, 8-13]
AI Orchestration & Platforms	AI Orchestration / Enterprise IT	Orchestration layers coordinate multiple models, agents, and data sources into unified workflows, reducing tool-call failures by	Transport latency (2-200ms), schema validation quality, deterministic vs. probabilistic control, and tool selection noise.	Risks of "tool bloat," shadow AI adoption, and 100% tool spoofing success rates in some protocols (MCP); sensitive data leakage (27.4% of AI inputs).	Transition to automated Context Platforms; use Directed Acyclic Graphs (DAGs) for task management; implement centralized planning with	[2, 11, 14-19]

		60-80% and manual work by 50%.			federated execution.	
LLMs in Electronic Health Records (EHR)	Biomedical / Health Informatics	Input structure (chunking) and model capacity are primary drivers of accuracy. DeepSeek-R1 achieved the highest performance (0.89 macro-F1).	Chunking strategy, aggregation rules (any-positive vs. majority), and phenotype difficulty (temporal logic/numeric inference).	Systematic overconfidence; temporal disambiguation failures; ontological granularity errors. Self-reported confidence is not a calibrated probability.	Prefer whole-document inference; avoid majority voting; adopt tiered decision policies (auto-accept vs. human review); use standardized templates.	[20]
Long-Context Evaluation Benchmarks	Machine Learning Evaluation	Traditional static QA is insufficient; new benchmarks (RULER, HELM, GSM-∞) reveal sharp degradation in complex tracing and aggregation as context length grows.	Reasoning complexity (operation count), signal-to-noise ratio, "indistinguishable noise" (Spider Topology), and "usable" vs. "advertised" length.	Benchmarks with low complexity fail to verify the need for expensive long-context models, leading to inefficient resource allocation.	Use Area Under Curve (AUC) for performance modeling; utilize synthetic data to minimize memorized knowledge; move beyond single-needle tests.	[6, 21-25]
Memory Systems in AI Agents	AI Systems Design / NLP	Memory is critical for long-horizon utility. Frameworks include internal (KV Cache, Weights) and external (Vector Index, Hierarchical Store) substrates.	Context explosion, access speed (internal) vs. scalability (external), and environmental complexity.	Persistent memory introduces privacy/security risks (targeted extraction attacks); "state drift" as complexity scales.	Use self-adaptive architectures that prune experience; implement Hierarchical Cognitive Caching (HCC); use vector databases for semantic search.	[21, 26, 27]
Reasoning Models (LRMs)	Legal Analysis / Deep Learning	Models (e.g., o1) use "chain-of-thought" to think step-by-step, providing logically consistent	5-10x higher cost/latency; "thinking token" sparsity (only ~20% of tokens are critical for steering the model).	Reliability is higher but models still hallucinate and lack true common sense; can follow	Be specific about reasoning steps; request "show your work" explanations; use a two-stage approach (fast	[1, 28]

		results for complex multi-factor analysis.		logical chains based on faulty premises.	models for filtering, reasoning models for depth).	
RAG & Retrieval-Exclusive Generation (REG)	Legal Investigations / AI Architecture	RAG retrieves passages as context; REG forces the AI to use *only* retrieved info. This dramatically reduces hallucinations but depends on retrieval quality.	Retrieval quality; search step accuracy; information acquisition bottleneck (uncontrolled retrieval leads to context waste).	Requires systematic verification and mandatory citations to maintain professional accountability.	Implement "Find-Analyze-Answer" frameworks; require citations for every substantive statement; use links for human verification.	[1, 29]
UX Design & Progressive Disclosure	UX Design / HCI	Techniques that present only necessary info to manage cognitive workload and reduce errors. Includes staged, conditional, and contextual disclosure.	Human working memory limits (Miller's findings); interaction cost; navigation depth (risk if > 3 levels).	Risk of "discoverability debt"; incorrect WAI-ARIA implementation can lead to accessibility failures or AT user fingerprinting.	Follow the "inverted pyramid" concept (critical info first); ensure interactive elements are keyboard focusable; align with WAI-ARIA standards.	[15, 30-34]
Technical Architectural Constraints (RoPE/BFloat16)	Transformer Architecture / NLP	Precision and embedding parameters (RoPE base values) define the absolute lower bound for context capability and numerical coherence.	Numerical precision (BFloat16 vs. Float32), RoPE base value, and out-of-distribution (OOD) rotation angles.	Inappropriate base values result in "superficial" long-context ability; BFloat16 causes cumulative errors starting from the first token.	Use "AnchorAttention" (fixed position ID for the first token); ensure the RoPE base is sufficiently large to maintain token discrimination.	[35, 36]

Virtual Patient Systems	Biomedical / Medical Education	LLM systems simulate internal medicine well but lack multimorbidity representation. Knowledge Graphs (KGs) improve accuracy by 16.02%.	Prompt design (few-shot), KG integration, and sample size consistency.	Hallucination rates (0.31%–5%) pose clinical risks; dataset bias (e.g., ICU-focused) limits generalizability to outpatient care.	Use hybrid KG-chain-of-thought architectures; aim for Top-1 accuracy ≥ 0.80 and System Usability Scale (SUS) ≥ 80.	[37]
Multimodal LLM Evaluation (MM-NIAH)	Multimodal LLM Evaluation	MLLMs perform significantly worse with image needles than text needles, struggling once context exceeds five images.	Context length (1k to 72k tokens), image resolution (patch count), and text-to-image ratio.	MLLMs fail to follow instructions and produce gibberish in very long multimodal contexts.	Improve multi-image perception; increase instruction-following ability; use specialized multimodal benchmarks.	[3]
Computational Cognitive Load (ICE Benchmark)	AI / Cognitive Science	Extraneous information (Saturation) and task-switching (Attentional Residue) drive performance degradation.	Extraneous load percentage (20%, 50%, 80%), procedural similarity, and intrinsic complexity (hop count).	Vulnerabilities to "cognitive-load overflow" can be weaponized in jailbreak attacks to compromise reasoning.	Use dynamic, load-aware stress testing (ICE benchmark); gate robustness assessments on baseline control settings.	[38]
AI Governance & Defense in Depth	Business Strategy / Cybersecurity	Security requires multi-layered mitigation across Model, Safety, Application, and Positioning layers.	Blast radius (scope control); input data quality; human-in-the-loop oversight.	Data security (72%) and legal implications (59%) are primary risks; agent hijacking and sensitive data leakage.	Design agents as microservices (bounded capabilities); implement "Least Permissions"; standardize documentation for use cases.	[39-41]

Gold Context Size in Aggregation	General / Biomedical / Math	Smaller gold context spans lead to worse performance and higher positional sensitivity compared to larger spans.	Gold context size (Small/Medium/Large), primacy effect, and distractor volume.	Aggregation failures degrade trust and factual correctness in downstream agentic tasks.	Strategically balance retrieved document sizes; implement structured evidence expansion.	[42]
YunChang: Unified Sequence Parallel	Deep Learning Infrastructure	A unified sequence parallel approach (USP) that combines DeepSpeed-Ulysses and Ring-Attention for training.	Number of attention heads and communication/computation overlap efficiency.	Ring-Attention poses risks of communication deadlocks in large-scale asynchronous deployments.	Apply hybrid parallelism (e.g., ulysses_degree=2, ring_degree=4); reorder tokens for load balance.	[43]
Extract-Then-Phenotype (ETP)	Biomedical / Clinical NLP	Two-step process (extraction then rule-based classification) that underperformed direct classification in this study.	Diagnostic importance weights of clinical elements and the learned threshold (λ).	Requires high-quality clinician consultation; failure to extract key entities leads to score miscalculation.	Weights should be aggregated at the record level; explicit definition of thresholds is required.	[20]
Lighthouse Attention	Deep Learning / NLP	A method developed for long context pre-training of models.	Context window size and attention mechanism efficiency.	Not in source	Use for handling extended sequence lengths during pre-training.	[8]
Orchard Agentic Modeling Framework	AI Orchestration	An open-source framework designed for agentic modeling and orchestration.	Framework scalability and interoperability with open-source models.	Not in source	Utilize for standardized agent development and orchestration logic.	[8]

[1] 4TH EDITION - Amazon S3

[2] Agent Harness for Large Language Model Agents: A Survey[v1] | Preprints.org

[3] Needle In A Multimodal Haystack - NIPS papers

[4] SwayZone Command Center Project Exo Operator Manual (1).pdf

[5] The Cognitive and Computational Limits of Information Processing: A Comprehensive Analysis of Context Saturation in Large Language Models

[6] Lost in the Middle: How Language Models Use Long Contexts

[7] Long-Context Models vs. RAG: When the 1M-Token Window Is the Wrong Tool - TianPan.co

[8] Daily Papers - Hugging Face

[9] 4 Enterprise AI Trends that will Define 2025 - Uniphore

[10] The Orchestration of Multi-Agent Systems: Architectures, Protocols, and Enterprise Adoption

[11] Beyond the Bot: Why Multi-Agent Orchestration is the Future of Enterprise AI - AppsTek Corp

[12] Autonomous AI agents 2026: the new rules for business governance - Raconteur

[13] Enterprise AI: From Copilots to Multi-Agent Ecosystems - Cognizant

[14] From RAG to Context - A 2025 year-end review of RAG - RAGFlow

[15] SwayZone Command Center Project Exo Operator Manual (1).md

[16] The AI Fragmentation Crisis: Orchestrating the Transition from Tool Bloat to Agentic Autonomy

[17] 10 AI Orchestration Platform Options Compared for 2026 - Domo

[18] 8 Ways AI Agents Are Evolving in 2026 - Salesforce

[19] Best AI Orchestration Tools in 2026: Enterprise Guide | Elementum AI

[20] A Systematic Exploration of LLM Behavior for EHR Phenotyping | medRxiv

[21] Rethinking Memory Mechanisms of Foundation Agents in the Second Half: A Survey - arXiv

[22] GSM-∞: How Do Your LLMs Behave over Infinitely Increasing Context Length and Reasoning Complexity? - arXiv

[23] RULER: What's the Real Context Size of Your Long-Context Language Models?

[24] HELM Long Context - Stanford CRFM

[25] Multi Needle in a Haystack - LangChain

[26] A Context-Aware Feedback Loop for AI-Assisted Verification IP Synthesis: Bridging the Gap from Natural Language to Regression-Ready - MDPI

[27] Toward Ultra-Long-Horizon Agentic Science: Cognitive Accumulation for Machine Learning Engineering - arXiv

[28] Dynamic Thinking-Token Selection for Efficient Reasoning in Large Reasoning Models

[29] SCALING SEARCH-AUGMENTED LLM REASONING VIA ADAPTIVE INFORMATION CONTROL - OpenReview

[30] Accessible Rich Internet Applications (WAI-ARIA) 1.3 - W3C on GitHub

[31] Progressive Disclosure - Everyday Concepts

[32] Progressive Disclosure Controls - Win32 apps | Microsoft Learn

[33] Progressive disclosure UX for responsive websites - Justinmind

[34] The Architecture of Progressive Disclosure in Hierarchical Information Systems

[35] When Precision Meets Position: BFloat16 Breaks Down RoPE in Long-Context Training

[36] NeurIPS Poster Base of RoPE Bounds Context Length

[37] Large Language Model–Based Virtual Patient Systems for History-Taking in Medical Education: Comprehensive Systematic Review

[38] arxiv.org

[39] Defense in depth for autonomous AI agents | Microsoft Security Blog

[40] AI Governance Best Practices | SAP LeanIX

[41] AI governance checklist: A basic framework for your business - Wipfli

[42] Lost in the Haystack: Smaller Needles are More Difficult for LLMs to ...

[43] YunChang: A Unified Sequence Parallel (USP) Attention for Long Context LLM Model Training and Inference. - GitHub